

SELECT among the model's own candidate futures

Best-of-4 search

make 4 candidate next chunks, keep the best one (128 clips)

Rewind

re-roll a chunk that froze, keep it if it moves more (32 clips)

EDIT the video while it is being generated

AdaSteer (our method from slide 6)

fine-tune the model on the video while generating (8 clips)

Noise warp

shift the input noise along the estimated motion (8 clips)

Fast-weight memory

write recent frames into extra trainable weights (8 clips)

