

SELECT among the model's own candidate futures

Best-of-4 search

make 4 candidate next chunks, keep the best one (128 clips)

Rewind

re-roll a chunk that froze, keep it if it moves more (32 clips)

EDIT the video while it is being generated

AdaSteer (our method from slide 6)

fine-tune the model on the video while generating (8 clips)

Noise warp

shift the input noise along the estimated motion (8 clips)

Fast-weight memory

write recent frames into extra trainable weights (8 clips)

Prediction slide

shift each predicted frame along the motion (8 clips)

Extra attention sink

keep more early frames in memory permanently (32 clips)

